

# An Exploration of Hierarchical Quilted Self-Organizing Maps

Theodore Hilk, Joseph Lynch

March 23, 2014

## Abstract

We explore the paper “Biomimetic sensory abstraction using hierarchical quilted self-organizing maps,” by J. W. Miller and P. H. Lommel of the Charles Stark Draper Laboratory. We place the paper in historical context, provide a high-level summary, and briefly review related literature. We give both a qualitative overview and a specific mathematical formulation of the hierarchical quilted self-organizing map (HQSOM) model, and provide a thorough, reproducible record of our design and implementation process for a Python-based HQSOM framework. We present our results from conducting each of the experiments described in Miller and Lommel’s paper, which involved the classification of various shapes invariant to shift and scale transformations in 3x3- and 7x7-pixel fields. We obtain results consistent with those of the authors, albeit noting considerable dependence on parameter values and the ambiguity of a certain key formula in the paper. Subsequently, we develop and test several improvements to the HQSOM algorithm itself. Some result in outperforming the original in classification tasks, especially when noise is introduced. We then implement preprocessing adaptations to allow us to apply both the original and improved HQSOM algorithms to audio, and we succeed in developing a genre classifier with excellent out-of-sample performance. We conclude by describing further possibilities for extension and discussing the broader significance of this line of research.

## 1 Introduction and Motivation

Biologically-motivated algorithms have a long history in artificial intelligence and machine learning. Indeed, AI as a field was born out of attempts to simulate human intelligence using artificial systems[1], and one of the earliest techniques, the perceptron, was explicitly based on a simplified model of a biological neuron.[2]

Although symbolic approaches enjoyed much more prominence following the publication of papers recognizing the limitations of early neural networks in the early 1960s[3], their own difficulties in dealing with problems of perception, learning, and pattern recognition led to a resurgence of interest in sub-symbolic systems from the 1980s onward.[1, 4, 5] It was during this period that the field of machine learning gained substantial prominence, focusing on pattern recognition, classification, and adaptive control, and placing a heavy emphasis

on quantitative and often statistical methods as opposed to symbolic and logical ones.[5]

Beginning largely in the early 1990s, collaboration in turn between machine learning researchers and neuroscientists produced models of neurological and cognitive systems demonstrating both strong concordance with empirical observations and excellent performance in the tasks addressed by those same systems.[6] Investigation in the area from that time onward has focused on the theory and simulation of increasingly substantial and complex aspects of human neuroscience, with a particular emphasis on neocortical functionality, the problem of vision, and the use of hierarchical, invariant representations for sensory stimuli.[6, 7] The significance of such modeling efforts can hardly be overstated: they have led to new findings in neuroscience, improvements in brain-computer interface technology, and perhaps most profoundly, a new paradigm in artificial intelligence research.[7, 8]

Currently, the most prominent of these models include the Neocognitron (Fukushima), HMAX (Riesenhuber, Poggio, Serre), the Neural Abstraction Pyramid (Behnke), Hierarchical-Temporal Memory (HTM) (George, Hawkins), Adaptive Resonance Theory (Carpenter, Grossberg), and VisNet (Stringer, Rolls).[7] With the exception of HTM, each of these approaches either relies on substantial hard-coded *a priori* knowledge specific to vision, neglects the role of temporal associations in learning, or requires separate training and execution phases in use.[7, 9] HTM in turn involves complex and inelegant models of “neurons”, though they bear little resemblance to those modeled by computational biologists, and at its core is nothing more than a recurrent spatio-temporal clustering algorithm incorporating interlayer feedback to provide auto-associative prediction capabilities.[9]

Our goals in undertaking this project were to identify and generalize the essential aspects of existing hierarchical cortical models, and to extend and improve upon techniques from the literature. Therefore, we chose the simplest and most general such model we could find that was not subject to the above-enumerated limitations. This was the Hierarchical Quilted Self-Organizing Map (HQSOM) model developed by Miller and Lommel, which we have succeeded in analyzing, reproducing, testing, improving, and extending.

Specifically, we have implemented the HQSOM model in numerical Python, conducted the same visual shape recognition experiments as the paper’s authors over a wide range of parameter values, and succeeded in reproducing their results in spite of ambiguity surrounding a particular key formula in the paper. Further, we have characterized the significance and implications of the HQSOM’s parameters, improved the noise performance of the algorithm by including a mean-squared-error-based activation function, and enhanced its separability characteristics by using a peripherally inhibitory “Mexican hat” neighborhood function instead of the standard gaussian. We have also improved its convergence rate without compromising representation quality, through the incorporation of an adaptive learning rate that is applied to the best-matching input prototype when a given input differs greatly from all current prototypical representations. We have also improved the efficiency of the training process by creating a “reset” function, which both eliminates the need for blank input sequences between successive training sequences, and ensures that no residual effects can carry over between these training sequences. Finally, we have ex-

tended our implementation to represent and classify audio, and successfully built a system to learn and classify music based on genre which performs very well out-of-sample. We believe that even greater potential exists for the improvement and generalization of this approach, and we highlight some of our views on the matter in our conclusion below.

## 2 Paper and Summary

We chose to replicate the paper “Biomimetic sensory abstraction using hierarchical quilted self-organizing maps,” by Jeffrey W. Miller and Peter H. Lommel of the Draper Laboratory. The paper develops a novel classification algorithm called a hierarchical quilted self-organizing map (HQSOM), then applies it to two simple visual processing tasks. The motivation for this approach lies in the desire to model the operation of the human visual cortex in a way that improves upon the realism of earlier techniques by: relying more on learning and less on hard-coded *a priori* knowledge, using temporal associations and not just spatial ones to construct representations, and learning on-line instead of requiring a separate training period.

The paper discusses several benefits of modeling the brain, noting that more accurate models have allowed for the creation of improved AI techniques, neurophysiological discoveries, and even better brain-computer interfaces. The authors choose to focus on vision in particular because it has such rich applications in software and is handled so easily by animals.

They then discuss various aspects of brain structure and function, noting that the isocortex, comprising about 85% of the brain in humans, is responsible most sensory processing, motor control, language, logic, mathematics, and spatial visualization. In spite of this, it appears to have a uniform, hierarchical structure throughout its various functional regions, and a substantial portion of the contemporary cognitive neuroscience community posits that the elements of this structure perform essentially the same information processing operations throughout the cortex.

Examining the ventral stream of the visual cortex in particular, the authors note that its regions are wired in a hierarchy. At the base, V1 neurons respond to small lines of various orientations in their respective and narrow portions of the visual field, while neurons in the higher V2 region respond to somewhat less localized simple shapes, those of the even higher V4 region respond to more complex shapes across larger areas, and those in the IT region respond invariantly to complex objects throughout the entire field of view. The authors then mention a few similar cortical models, noting the above limitations in realism. Such alternative models are described in more detail in our literature review below.

Miller and Lommel define the goals for their own model in terms of biological functional equivalence, generalizability across multiple sensory domains, and simplicity. They specifically note that computational efficiency was only a secondary goal. In pursuing functional equivalence, the authors take the explicit position that there exists some basic “building block” or unit of the cortex beyond which no further details must be modeled, so long as the abstract func-

tionality of the unit is properly preserved.

They describe this functionality in terms of the unsupervised learning of abstract patterns in the input data, and of the recurrent learning of patterns in the sequences of inputs patterns at lower levels of the hierarchy. They propose that this reduction of sensory input to abstract concepts may be conceived in terms of the related processes of spatial and temporal clustering, or "pooling". Spatial pooling learns patterns of input data that are spatially co-incident at a given level of the hierarchy, while temporal pooling learns patterns that tend to occur near one another in time.

The authors note that numerous algorithms exist for spatial clustering, including K-means, SOMs, expectation maximization, winner-take-all neural networks. This is discussed further below in terms of our current and future extensions to the HQSOM model. Likewise, temporal clustering and sequence processing can be performed with anything from basic statistical techniques like ARMA, ARIMA, NARMA, GARCH, N-GARCH, etc. to recurrent neural networks, temporal Kohonen maps (TKMs), or the simpler and faster-converging recurrent self-organizing maps (RSOMs).

In keeping with this approach, the HQSOM algorithm creates invariant, low-dimensional representations of high-dimensional data by passing them through successive stages of spatial and temporal pooling. Spatial pooling is accomplished by means of self-organizing maps (SOMs), while temporal pooling is handled by a type of modified SOM called a recurrent self-organizing map (RSOM). The authors elected to use the SOM model because it was familiar to them, and because it is relatively simple and well-researched. Likewise, they chose RSOMs for temporal pooling due to their demonstrated performance in the area and their elegant similarity to the ordinary SOM.

A SOM, or Kohonen network, is a type of unsupervised nonlinear dimensionality reduction algorithm, often used for data visualization. It can also be viewed as a form of vector quantization. SOMs operate by building a table of "weight vectors" in the input space, then modifying the weight vectors closest to each successive input vector to make them closer to that input vector, such that the set of weight vectors overall converges to yield a summary representation of the clusters in the input data in a topology-preserving fashion. An RSOM is a SOM that takes an exponential moving average over the differences between input vectors and weight vectors when computing these distances, which allows measures of similarity to extend across a succession of inputs and hence into the temporal domain. Both SOMs and RSOMs are described in more mathematical detail in the section entitled *System Design and Variables* below.

After describing the motivation and technical details of HQSOMs, Miller and Lommel's paper describes the recursive combination of HQSOM base units, which are comprised of a SOM whose regularized activation vector over its weight vectors is provided as input to an RSOM (see below for specific mathematical details). The SOM performs spatial clustering over its input, yielding a set of representative spatial feature vectors in the input space, and the RSOM proceeds to perform temporal clustering over sequences of these spatial patterns, such that the RSOM's weight vectors each correspond to some particular spatiotemporal cluster in the input data. Thus, the HQSOM base unit as a whole produces a dimensionally-reduced representations of the spatiotemporal

contents of its input. It produces spatiotemporal abstractions. Composing HQSOMs in a hierarchical arrangement, such that the activation vector over the weight vectors of the RSOM in a given HQSOM base unit are provided as the inputs to the SOM of another HQSOM base unit, allows for the representation of progressively more invariant spatiotemporal patterns in the input data.

The authors then describe two experiments involving the learning and classification of simple visual patterns, which we describe in detail and reproduce below.

### 3 Hypothesis, Plan, and Risks

We wish to begin by re-implementing the HQSOM algorithm described in the paper, using numerical Python. The primary risk in doing so is that the algorithm is rather complex, and certain aspects are not thoroughly described. We also expect that it will take a long time to code, and that the high dimensionality and fairly opaque nature of the HQSOM's representations may make it difficult to debug our implementation. For example, if a network is failing to produce the desired result, looking at the RSOM unit map is not particularly helpful.

We will then attempt to reproduce the experimental results presented by Miller and Lommel, which showed that HQSOMs could form shift- and scale-invariant representations of various shapes within 3x3- and 7x7-pixel fields of vision. We expect that the testing framework and test data may take a considerable amount of time to create. We also anticipate that the experimental results may be difficult to reproduce due to parameter sensitivity, and that computational efficiency may pose hurdles given the number of training cycles used in the paper. As above, we again note that it can be difficult to debug these types of networks due to the internal structure being fairly mathematical in nature. Finally, the potential fragility of these types of systems is well known in the academic community, so it is very possible that we will not be able to get any viable results whatsoever due to bad parameters and long test runs.

Once we are able to replicate these results to a reasonable degree of precision, we will extend the HQSOM framework into the audio domain and attempt to classify music by genre. The main risks emerge from the potential complexity of these extensions, the nuances involved in producing high-quality spectrograms, and the anticipated large size and computational burden of networks capable of classifying something as abstract as genre.

### 4 System Design and Variables

HQSOM networks are comprised of hierarchically stacked building blocks known as HQSOM base units. An HQSOM base unit in turn consists of a stacked SOM-RSOM pair, such that input to the base unit as a whole is provided as the input to the SOM, then the regularized activation vector over the map space of the SOM is provided as the input to the RSOM, and finally the activation vector over the map space of the RSOM is provided as the input to any HQSOM base units stacked above this one. If this is the top unit in the network, then this

final RSOM activation vector yields an invariant summary of the input data as a whole and may be used for classification. Composing these SOM-RSOM pairs into multi-layered networks yields Hierarchical Quilted Self-Organizing Maps (HQSOMs), which can identify both spatial and temporal clustering in data over multiple levels of abstraction. The general use case for HQSOMs is to identify spatiotemporal clusters in input data, such that a supervised learning technique can be applied to actually make the classification, since the clusters themselves rely on some associated semantic meaning in order to be regarded as labels. However, for the sake of this paper, we will take such semantic meaning for granted and thus consider clustering equivalent to classification. The following discussion of SOMs, RSOMs and HQSOMs is based on the corresponding sections of the Miller and Lommel paper, and hence the vast majority of this method explanation can be found in that paper as well [7].

#### 4.1 SOM

The basic SOM computational block can either be trained on data or asked to classify data. The SOM is made up of a  $m \times n$  matrix that maps inputs of dimension  $m$  to outputs of dimension  $n$ . For example a SOM designed to take in 3d vectors and output a 5d vector looks like:

$$\begin{pmatrix} .3 & .7 & .1 & .14 & .01 \\ .3 & .1 & .01 & .16 & .9 \\ .3 & .03 & .8 & .7 & .01 \end{pmatrix}$$

Each column is a map unit, and whichever map unit  $\mathbf{w}_b$  is closest to the input  $\mathbf{x}$  is considered the best-matching unit (BMU). The measure of closeness is usually simply Euclidean distance, so:

$$\mathbf{w}_b = \underset{w_i}{\operatorname{argmin}} \|\mathbf{x} - \mathbf{w}_i\| \quad (1)$$

During the training stage, input vectors are applied to the input and then an update rule is applied over the entire map space that shifts map units that are nearest to the input data towards the input data:

$$\mathbf{w}_i(t+1) = \mathbf{w}_i(t) + \gamma h_{ib}(t)(\mathbf{x}(t) - \mathbf{w}_i(t)) \quad (2)$$

where gamma is the rate of learning,  $h_{ib}$  is the neighborhood function, and  $\mathbf{w}_i$  is the map unit being modified. The neighborhood function is defined as a function that is close to zero for units far away from the BMU. Traditionally a Gaussian is used:

$$h_{ib}(t) = \exp\left(\frac{-\|I_i - I_b\|^2}{\mu(t)\sigma^2}\right) \quad (3)$$

where  $I_i$  indicates the index of the  $i$ th unit,  $\mu(t)$  is a decreasing function of mean squared error and  $\sigma$  is the learning radius. A SOM therefore has two parameters that need to be tuned: the learning rate  $\gamma$  and the learning radius  $\sigma$ . For example, two sample unit maps after an update with  $\mathbf{x} = (.1, .1, .1)$ ,  $\gamma = .2$  and different sigmas would look like: Clearly, the update with the larger sigma affected more of the map space. Also it is important to note that units were

$$\begin{pmatrix} 0.26 & 0.7 & 0.1 & 0.14 & 0.01 \\ 0.26 & 0.1 & 0.01 & 0.16 & 0.9 \\ 0.26 & 0.03 & 0.8 & 0.7 & 0.01 \end{pmatrix}$$

Figure 1:  $\sigma = 1$

$$\begin{pmatrix} 0.26 & 0.607 & 0.1 & 0.139 & 0.01 \\ 0.26 & 0.1 & 0.017 & 0.159 & 0.897 \\ 0.26 & 0.041 & 0.748 & 0.687 & 0.01 \end{pmatrix}$$

Figure 2:  $\sigma = 100$

pulled towards the input, but with a less dramatic effect as the map index increased (separating itself further from the BMU).

During activation, the SOM can return two types of activation vectors:

1. Discrete: A vector of the correct dimension with the BMU index set to 1 and all others set to 0.
2. Continuous: A vector  $A(t)$  defined as the normalized form of a vector constructed as follows:

$$A_i = \frac{1}{\|\mathbf{x}(\mathbf{t}) - \mathbf{w}_i\|^2} \quad (4)$$

## 4.2 RSOM

The Recurrent SOM is an extension of the basic SOM that adds an exponential moving average of differences between observed inputs and units in the map with time-decay parameter  $\alpha$ . At each update the differences are updated and instead of looking for the BMU in map space, the BMU index is chosen by finding the minimum magnitude recursive difference. Furthermore, instead of applying  $\mathbf{x}(t)$  directly to the map, the recursive difference for a particular unit is applied in each unit's update rule and  $\mathbf{x}(t)$  is used to update the recursive difference matrix:

$$\mathbf{y}_i(t+1) = (1 - \alpha)\mathbf{y}_i(t) + \alpha(\mathbf{x}(t) - \mathbf{w}_i(t)) \quad (5)$$

The update rule becomes:

$$\mathbf{w}_i(t+1) = \mathbf{w}_i(t) + \gamma h_{ib_r}(t)\mathbf{y}(\mathbf{t}) \quad (6)$$

where the neighborhood function is computed using the recursive BMU as the BMU index instead of the map space BMU. The proper tuning of  $\alpha$  depends on how responsive one wishes for the RSOM to be: lower values of  $\alpha$  cause the moving average of inputs to dominate (long term memory), whereas higher values of  $\alpha$  (close to 1) mean that recent inputs dominate (short term memory)

## 4.3 SOM-RSOM Pair

The final computational structure is the SOM-RSOM pair. Recognizing that SOMs are effective at identifying spatial clustering, and RSOMs are effective

at identifying temporal clustering, but SOMs do zero temporal clustering and RSOMs have degraded spatial clustering, the SOM-RSOM pair is intended to get the best effect of both. First the input data is fed into the SOM to get a spatial representation, and then this representation is fed into the RSOM to do temporal clustering.

## 5 Implementation Process and Results

Two implementations of SOMs and RSOMs were written, one being as close as possible to the reference implementation for reproducibility, and the other having a number of slight changes to improve convergence. As is to be expected, the reference paper leaves out a number of implementation details so this is our best approximation.

### 5.1 Reference Implementation of SOM and RSOM Units

Self-Organizing Maps were implemented via a Python class with three main methods: a constructor, an update method that takes a numpy array representing the input vector as well as the learning parameters (gamma, sigma, etc...) to use for that particular update and modifies the internal state of the SOM accordingly, and a method to request the activation vector for a given input vector. Internally the SOM map was stored as a  $m \times n$  numpy array where  $n$  is the input vector size and  $m$  is the size of the map space. During an update call, the Best Matching Map Unit (BMU) for any given input vector was determined using a linear search for the minimum Euclidean distance and then all map units near to the BMU as well as the BMU were shifted towards the input according to a Gaussian neighborhood function. The standard SOM update rule was used as per equation (2).

A linear search was preferred due to the high dimensionality of the space, and a Gaussian neighborhood function was chosen for reproduceability. The activation method returned either a discrete or a continuous representation of the map's activation to a given input. The discrete representation is defined as above, and the normal representation is given in equation (7) where  $a_i$  represents the  $i$ th position in the activation vector,  $\mathbf{w}_b$  is the BMU,  $\mathbf{x}$  is the input vector, and  $\mathbf{w}_i$  represents the  $i$ th map unit.

$$a_i = \frac{1}{\|\mathbf{w}_i - \mathbf{x}\|}, \mathbf{a} = \frac{\mathbf{a}}{\|\mathbf{a}\|} \quad (7)$$

Recurrent Self-Organizing Maps were simply a subclass of the SOM that use the modified RSOM update rule as well as storing the recursive difference matrix as a numpy array. The time decay parameter ( $\alpha$ ) was passed in at every update call.



## 5.2 Basic Design of SOM-RSOM Pair and Other Hierarchical Structures

Since both SOMs and RSOMs are implemented as Python objects, the SOM-RSOM pair simply consisted of a SOM object and a RSOM object with an update and activation method that takes in an input vector  $\mathbf{x}$ , feeds it into the SOM to get a transformed activation vector  $\mathbf{y}$  and finally takes that  $\mathbf{y}$  and feeds it into the RSOM to get the final output which is the BMU of the RSOM. The only difference between the SOM-RSOM update and activation methods is that the update method calls update internally (thus changing the state of the network), whereas the activation method merely passes along activation vectors. In the code this SOM-RSOM pair was referred to as a HQSOM because a single SOM-RSOM pair does indeed form the simplest HQSOM.

Hierarchies were built at first by hard wiring these HQSOM base units together. However, in order to facilitate testing of the audio extension, a framework was designed that allowed for any arbitrary tree HQSOM accepting input with a 1-dimensional topology (e.g. a line of pixels or spectral power densities, as opposed to a 2d image). The first level of the tree reads data from the input, and passes its activation vector to the next layer, which passes its activation vector to the next layer, etc ... The output of the top level node is the representation of the input that is (hopefully) invariant under certain conditions.

## 5.3 Replication of First Experiment

The first experiment presented in the paper was a simple example of 3x3 images with 3 pixel horizontal and vertical lines that have been shifted to all possible positions. This data set is small enough to be enumerated, and simple enough in concept to use a single SOM-RSOM pair as the HQSOM network. We implemented the network as a single HQSOM that had an input size of 9, internal SOM map size of 18, and internal RSOM map size of 3. We mapped the 3x3 image grids to a linear vector of size 9 by iterating through the image left to right and top to bottom. The mapping used is shown in Figure 3.

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{pmatrix} \rightarrow (1 \ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \ 8 \ 9)$$

Figure 3: Experiment 1 Image to Vector Mapping

The implementation was shown to be correct by two different tests: performance on non-noisy data, and aggregate performance over many noisy data sets. During training, the HQSOM is exposed to three blank images, followed by three line images, followed by three blank images where the three line images alternate between the three horizontal and the three vertical images. An example training sequence (without noise) is shown in Figure 4. Applying the sequence shown in Figure 4 hundreds of times with parameters  $\gamma_{som} = \gamma_{rsom} = .1$ ,  $\sigma_{som} = 16$ ,  $\sigma_{rsom} = 90$ , and  $\alpha = .1$  trained the HQSOM and clustered the weight vectors in the map units of the SOM and RSOM. Since the blank images are solely meant to reset the RSOM EMA difference matrix,



Figure 4: Experiment 1 Training Sequence

a method was added to RSOMs and HQSOMs that allows the difference matrix to be cleared (with a random index selected to have a value of .01 so that the BMU varies randomly after each reset) and the training steps with blank images were replaced with a call to this function. After roughly 2500 training samples were shown to the network, the HQSOM was asked to classify (return the BMU of the top level RSOM) each piece of data again and as expected all horizontal lines were classified the same, all vertical Lines were classified the same, and all blank images were classified the same regardless of position. The network had successfully formed an invariant representation of vertical and horizontal lines in a 3x3 field of view.

To test the noise tolerance of the network, Gaussian noise with standard deviation .1, .2, .25 and .3 was applied to the test data, which was then trained on as before and the HQSOM was again asked to classify samples of noisy vertical and horizontal lines (with different noise from the training data). The results of 100 trials with each noise level are summarized in the following table:

Noise Std. Deviation	Number Correctly Clustered
.1	99/100
.2	31/100
.25	5/100
.3	4/100

Clustering “correctly” simply means that all vertical lines had the same BMU at the output, all horizontal lines had the same BMU at the output but different from the vertical lines, and all blank images had the same BMU at the output but different from either the vertical or horizontal lines. For example: (1 0 0 0 2 2 2) is a “correct” clustering if we apply one blank image, three vertical lines, and three horizontal lines but (2 0 0 0 0 0 0) is not. It is worth noting that when fewer than 2000 training steps were taken the map would often converge to a “Something vs Nothing” map in that the HQSOM would be very good at clustering lines together and blank images together, but would not differentiate between the two types of lines. This makes sense because the alpha is small enough that it could take time for the final representation to form.

## 5.4 Replication of Second Experiment

The second experiment presented in the paper aimed to create shift and scale invariant representations of squares, diamonds, and X shapes in a 7x7 grid. To replicate this experiment, a two tiered network was created such that there were 9 low level HQSOMs that each inspected a 3x3 swatch of the 7x7 grid with 1 pixel overlaps on each side. These base units fed into a top level HQSOM having a 9 dimensional vector input (composed of the BMUs of each of the bottom level

HQSOMs), which in turn outputted a BMU index representing the cluster an input belongs to. The goal was to find a shift invariant representation of these shapes by exposing the network to each family of shapes that are both scaled and then shifted around in a spiral fashion, followed by blanks for 100 steps, and then the next shape, more blanks, and so on. The input sequence is shown in Figure 5.




Figure 5: Experiment 2 Input Data

The paper claims upward of 95% clustering, but we were unable to even achieve a four classifier using the specified parameters from the paper as per Figure 6.

	$\gamma$	$\alpha$	$\sigma$	Map Size
Layer 1 SOMs	.1	1	4	65
Layer 1 RSOMs	.01	.1	10	17
Layer 2 SOMs	.1	1	2	513
Layer 2 RSOMs	.001	.01	50	17

Figure 6: Experiment 2 Parameters

The best run of our HQSOM yielded the following final distribution where the final distribution is simply the number of test images classified as that BMU divided by the total number of test images for a given data set.

```
#####
Data Set: BLANK
Most Frequently Classified As (MODE): 4
Full Distribution over Final RSOM Map Space:
[ 0.  0.  0.  0.  1.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.]
#####
Data Set: SQUARE
Most Frequently Classified As (MODE): 3
Full Distribution over Final RSOM Map Space:
[ 0.          0.08571429  0.17142857  0.45714286  0.28571429  0.          0.
  0.          0.          0.          0.          0.          0.          0.
  0.          0.          0.          ]
#####
Data Set: DIAMOND
Most Frequently Classified As (MODE): 2
Full Distribution over Final RSOM Map Space:
[ 0.          0.14285714  0.77142857  0.08571429  0.          0.          0.
  0.          0.          0.          0.          0.          0.          0.]
```

```

0.      0.      0.      ]
#####
Data Set: X
Most Frequently Classified As (MODE): 2
Full Distribution over Final RSOM Map Space:
[ 0.      0.      0.97142857  0.02857143  0.      0.      0.
  0.      0.      0.      0.      0.      0.      0.
  0.      0.      0.      ]

```

While there is certainly convergence, it is not 95%. Due to the fact that runs of this particular simulation took well over 9 hours, we were only able to test about 10 different parameter combinations, none of which yielded a better result than the one supplied above. It appeared that the HQSOM was not using the full breadth of the RSOM map space in any of the HQSOM units, which may have resulted from an incorrect formulation of the activation vectors that were being passed up. We suspect this as the cause because the formulas presented in the paper for the continuous version of the activation vectors were not well-formed. Thus, we had to deduce the proper formulas based on somewhat unclear explanations of these vectors and their significance. It was also noticed that the networks frequently converged to certain states very quickly and then never moved from that state. The lack of a good activation vector and unintentional fast convergence led us to the following innovations:

1. A regularized activation vector based on the Mean Squared Error of the BMU vector.
2. The notion of an adaptive  $\gamma$ .
3. The use of a Mexican Hat neighborhood function (second derivative of Gaussian) instead of a Gaussian neighborhood function.

## 5.5 Changes to Algorithm and Relative Performance

Having completed the implementation of the paper's networks as far as deemed feasible, the three innovations previously mentioned were implemented in code and tested to compare relative performance. In all cases our implementation proved to be superior, especially when we began testing our Audio extension.

We began by implementing the new activation vector. Since the adaptive  $\sigma$  function was based on minimizing the mean-squared error (MSE) of the BMU and input vector, we believed that it would also be a good metric to measure activation by. This led to the definition in equation (8) where the vector parameters are the same as in equation (7) except that now  $MSE(\mathbf{x}, \mathbf{y})$  indicates the mean squared error between vectors  $\mathbf{x}$  and  $\mathbf{y}$ .

$$a_i = \frac{MSE(\mathbf{w}_b, \mathbf{x})^3}{MSE(\mathbf{w}_i, \mathbf{x})^3} \quad (8)$$

The errors were cubed so that the distribution in the vector would be more concentrated on the BMU's index.

Next we implemented the adaptive  $\gamma$ . During testing, the MSE between the input and BMU would often spike due to a training example that had never been seen before, which would cause  $\sigma$  to spike, which would in turn pull the entire map space drastically towards the new training example. This led to things like oscillations between two clusters while entirely losing a third or fourth possible clustering. To remedy this problem, we kept an exponential moving average of MSE in the SOMs of each HQSOM unit and whenever the MSE spiked by more than an order of magnitude, we made  $\gamma$  equal to some high fraction (in our case .6) just for the update of the BMU weight. This meant that a single map unit was pulled into the new cluster, becoming the semi-permanent BMU for all new training data that fell in that cluster.

Lastly we replaced equation 3 with equation 9.

$$h_{ib}(t) = (1 - \frac{||I_i - I_b||^2}{\sigma^2}) * \exp(\frac{-||I_i - I_b||^2}{\mu(t)\sigma^2}) \quad (9)$$

The advantage of this function is that it pushes away map units that are near to the input, but far enough away to be considered not matching. This helps deal with the pre-mature convergence noticed during testing. The negative aspect of this function is that the  $\sigma$  parameter has to be chosen very carefully so as to prevent uniform distributions from forming.

With our improvements in place, we re-ran the noise tolerance test from Experiment 1 and the large network test from Experiment 2. The results for the noise tolerance test over 100 runs is summarized in table 7.

Noise Std. Deviation	Paper Implementation	Our Implementation
.1	99/100	100/100
.2	31/100	73/100
.25	5/100	39/100
.3	4/100	12/100

Figure 7: Experiment 1 Test with Noisy Data: Clustering Results

Clearly our implementation represents a significant improvement in the face of noisy data. This is most likely due to the fact that we better allow for late cycle plasticity and are therefore able to compensate for early noisy examples with later less noisy examples.

Once again the second experiment proved challenging. We were unable to get any better results with the parameters provided in the paper, but by using the following network parameters, we were able to get a nearly equivalent result with far fewer cpu cycles:

	$\gamma$	$\alpha$	$\sigma$	Map Size
Layer 1 SOMs	.1	1	$\sqrt{20}$	40
Layer 1 RSOMs	.01	.1	$\sqrt{150}$	25
Layer 2 SOMs	.1	1	$\sqrt{15}$	150
Layer 2 RSOMs	.05	.02	50	7

```
#####
Data Set: BLANK
Most Frequently Classified As (MODE): 2
Full Distribution over Final RSOM Map Space:
[ 0.  0.  1.  0.  0.  0.  0.]
#####
Data Set: SQUARE
Most Frequently Classified As (MODE): 0
Full Distribution over Final RSOM Map Space:
[ 0.343  0.      0.314  0.171  0.143  0.029  0.   ]
#####
Data Set: DIAMOND
Most Frequently Classified As (MODE): 2
Full Distribution over Final RSOM Map Space:
[ 0.      0.      0.457  0.171  0.029  0.029  0.314]
#####
Data Set: X
Most Frequently Classified As (MODE): 6
Full Distribution over Final RSOM Map Space:
[ 0.      0.      0.2    0.2    0.2    0.057  0.343]
SUCCESS
```

When run with the same network, the reference implementation only produced two main statistical clusters. Once again, our implementation seemed superior.

## 5.6 Extension into Audio

The final stage of our project was extending our framework into audio classification. The goal was to be able to give a HQSOM spectrograms, and have the network cluster similar genres together. The first step of implementation was to build a generic framework that would allow for any arbitrary tree type network to be constructed without hard-wiring the HQSOMs together. Then, a test framework was created that allowed us to take 15 second snippets of songs, compute FFTs over windows that lasted  $\frac{1}{10}$  of a second for each song, and puncture the FFTs such that we reduced the input space to a 128 dimension vector. At the end of data processing we had the six spectrograms shown in Figure 8. The data were derived from the following songs:

1. Techno (Training) - Rudenko - Everybody
2. Techno (Testing) - Kid Cudi - Day and Night
3. Rock (Training) - Red Hot Chili Peppers - Californication
4. Rock (Testing) - Red Hot Chili Peppers - By the Way

5. Classical(Training) - Unknown Orchestra Conducted By George Winston  
- Carol Of The Bells
6. Classical(Testing) - Beethoven - Symphony No. 9




Figure 8: Spectrograms of Input Data For Audio Classification

The “Training” tag indicates that the network was exposed to sequential samples of that song’s FFTs in entirety, and then shown random 1 second clips from the spectrogram in rotation with the other Training songs. “Testing” songs were never exposed to the network during the training phase, but during the classification stage the HQSOM was asked to classify them. A network was built that consisted of two base SOM-RSOM nodes that take in 64 inputs each (half of the 128) and each output a BMU to a second SOM-RSOM, which then output their BMUs to a combined 2d vector which is used as the input to a final SOM-RSOM node which outputs the classification BMU (See the attached code for exact parameter details). Unfortunately, the reference SOM-RSOM implementation was entirely unable to separate the data, resulting in a 1-classifier for this network and every other network that we could conceive of. When we used this network in conjunction with our improvements, however, we were able to get the following positive results:

```
#####
Results for Techno
Final Distribution Over Map Space
[ 0.265  0.349  0.      0.386  0.   ]
MODE: 3
#####
Results for TechnoTEST
Final Distribution Over Map Space
[ 0.287  0.275  0.      0.438  0.   ]
MODE: 3
#####
Results for Classical
Final Distribution Over Map Space
```

```

[ 0.526 0.154 0.      0.321 0.   ]
MODE: 0
#####
Results for ClassicalTEST
Final Distribution Over Map Space
[ 0.49  0.202 0.      0.308 0.   ]
MODE: 0
#####
Results for Rock
Final Distribution Over Map Space
[ 0.434 0.554 0.      0.012 0.   ]
MODE: 1
#####
Results for RockTEST
Final Distribution Over Map Space
[ 0.266 0.734 0.      0.      0.   ]
MODE: 1

```

These are the results after cycling through each training song in its entirety, followed by three random 1-second window exposures for each song, making sure to clear the HQSOM's difference matrix when switching between songs. From them, we observe that the network has successfully separated the three different songs presented in the training data. More impressively, however, it has also successfully classified the out-of-sample data by genre, even though these data were not even taken from the same songs as those in the training set.

We note that the differences between the activation levels of the BMUs for different genres in this result were in some cases small, and thus that the network was close to misclassifying some of the data. We also note that it took quite a bit of parameter tweaking to obtain the above results. Given our limited time and computational resources, however, we consider this acceptable as a proof-of-concept.

Using only a single training cycle, we have successfully applied an HQSOM network to form invariant representations of the genres of three musical pieces. Further, we have applied these representations to correctly classify three additional songs to which the network had never been previously exposed. We did this with no hard-coded *a priori* knowledge whatsoever – the network obtained the sum total of its musical knowledge through exposure to a single play of each of the respective training songs.

## 6 Discussion and Conclusions

HQSOMs represent a promising path towards invariant spatio-temporal reasoning through massively parallel vector math. This paper has successfully reproduced many of the results of Miller and Lommel, improved upon the convergence and noise tolerance properties of their algorithm, and extended the HQSOM model into the audio domain where these types of networks show great promise. We have demonstrated that HQSOMs can simultaneously perform spatial and temporal clustering at multiple layers of abstraction, permitting invariant feature representation and classification over both space and time. Further, we have shown that they can do so in the presence of considerable noise, and that



they perform quite robustly out of sample – in particular, they can successfully classify never-before-seen audio data according to the highly abstract criterion of genre. All of this is done in a fully unsupervised fashion.

In performing our experiments, however, we have also identified a number of shortcomings in both the original HQSOM model that persist in our improved version. We have found that the networks are highly sensitive to parameter values and to some extent initial conditions, and that they often exhibit problems with premature convergence. We believe these issues present promising avenues for further research, with a particular focus on automatically tuning system parameters based on descriptive metrics for the input data and the sets of SOM and RSOM map units. This would effectively reduce the number of parameters requiring specification in the model, and would make network performance much less reliant on the empirical testing of configuration settings. As an example, the adaptive  $\gamma$  as shown in this paper is a good start, but these networks must become less fragile if they are to be widely adopted.

More broadly, our work has enabled us to characterize the essential properties of the general class of hierarchical, invariant, spatiotemporal representation and classification algorithms. Namely, they must perform recurrent spatial and temporal clustering over multiple levels of a hierarchy while preserving the underlying topology of their high-dimensional input data. It has not escaped our notice that SOMs are merely one of many topology-preserving techniques for nonlinear dimensionality reduction, and we strongly believe that the use of others such as locally linear embedding or Isomap in a sort of generalized spatiotemporal representation algorithm could yield promising results.

The incorporation of feedback mechanisms sending information from higher regions to lower ones also appears as an immediate possibility for extension, and including such capabilities would provide a means for auto-associativity and prediction much like that attempted by the HTM model.[9, 8] Having already incorporated the property of invariance, there is considerable reason to believe that a generalized spatiotemporal representation algorithm with such feedback could add emergence, reification, and multistability to its repertoire, completing its implementation of the core principles of gestalt systems.[9, 8] Much more distant possibilities include the incorporation of reinforcement learning, models of motor control and attention, and the capacity for episodic memory formation, with the eventual goal of constructing simulated or robotic intelligent agents that exhibit goal-directed behaviors and learn from their interactions with the world.[8] For now, however, we are pleased with our results and look forward to further investigation.

## References

- [1] M. Lungarella, F. Iida, J. C. Bongard, R. Pfeifer. AI IN THE 21ST CENTURY – WITH HISTORICAL REFLECTIONS. 50 Years of AI, Festschrift, LNAI 4850, pp. 18, 2007.
- [2] F. Rosenblatt. THE PERCEPTRON: A PROBABILISTIC MODEL FOR INFORMATION STORAGE AND ORGANIZATION IN THE BRAIN. Cornell Aeronautical Laboratory, Psychological Review, v65, No. 6, pp. 386408, 1958.

- [3] H. L. Dreyfus, S. E. Dreyfus. MAKING A MIND VERSUS MODELING THE BRAIN: ARTIFICIAL INTELLIGENCE BACK AT A BRANCHPOINT. *Daedalus* , v117, No. 1, Artificial Intelligence pp. 1543, Winter, 1988.
- [4] J. Schmidhuber. AI IN THE 21ST CENTURY – WITH HISTORICAL REFLECTIONS. 50 Years of AI, Festschrift, LNAI 4850, pp. 2941, 2007.
- [5] J. G. Carbonell, R. S. Michalski, T. M. Mitchell. MACHINE LEARNING: A HISTORICAL AND METHODOLOGICAL ANALYSIS. The AI Magazine, pp. 6979, Fall 1983.
- [6] T. Serre, M. Kouh, C. Cadieu, U. Knoblich, G. Kreiman, T. Poggio. A THEORY OF OBJECT RECOGNITION: COMPUTATIONS AND CIRCUITS IN THE FEEDFORWARD PATH OF THE VENTRAL STREAM IN PRIMATE VISUAL CORTEX. AI Memo 2005-036. CBCL Memo 259. Massachusetts Institute of Technology, Center for Biological and Computational Learning. Dec. 2005.
- [7] J. W. Miller and P. H. Lommel. BIOMIMETIC SENSORY ABSTRACTION USING HIERARCHICAL QUILTED SELF-ORGANIZING MAPS. The Charles Stark Draper Laboratory, Inc. 555 Technology Square, Cambridge, MA 02139-3563, USA. 2006.
- [8] J. Hawkins and S. Blakeslee. ON INTELLIGENCE. Tomes Books, Holt, New York, USA. 2002.
- [9] J. Hawkins, S. Ahmad, D. Dubinsky. HIERARCHICAL TEMPORAL MEMORY INCLUDING HTM CORTICAL LEARNING ALGORITHMS. Numenta, Inc. 811 Hamilton St., Redwood City, CA 94063, USA. Sept. 2011.